

# INGI2145: CLOUD COMPUTING (Fall 2015)

Beyond MapReduce: Higher-level languages, Graphs

22 October 2015

# MapReduce: Not for Every Task

- MapReduce greatly simplified large-scale data analysis on unreliable clusters of computers
  - Brought together many traditional CS principles
    - functional primitives; master/slave; replication for fault tolerance
  - Hadoop adopted by many companies
  - Affordable large-scale batch processing for the masses

But increasingly people wanted more!

# MapReduce: Not for Every Task

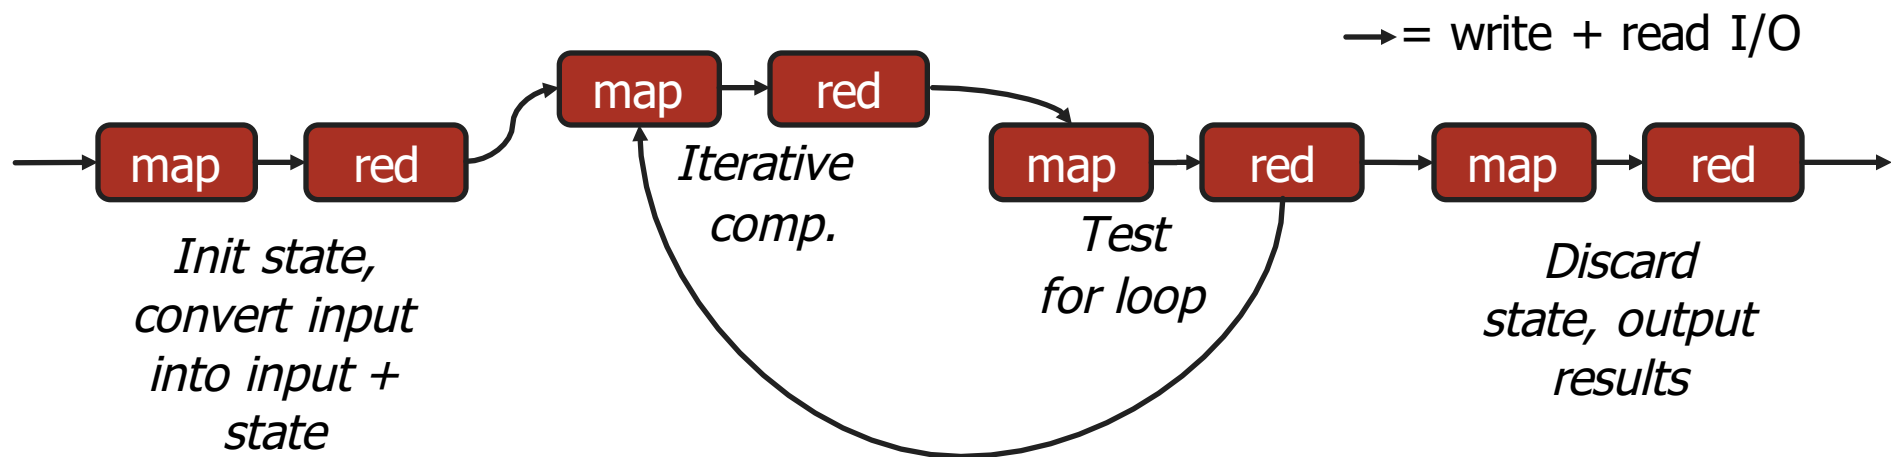
But increasingly people wanted more:

- More complex, multi-stage applications
- More interactive ad-hoc queries
- Process live data at high throughput and low latency

Which are not a good fit for MapReduce...

# MapReduce for Iterative Computation



- MapReduce is essentially functional
- Expressing iterative algorithms as chains of Map/Reduce requires passing the entire state and doing a lot of network and disk I/O
  - Recall all between-stage results are materialized to reliable and distributed storage (HDFS)



# MapReduce for Ad-hoc Queries

- MapReduce specifically designed for batch operations over large amounts of data
- New analysis task means writing a new MapReduce program
  - Tedious thing to do with languages such as Java
  - Programming interface is not familiar to traditional data analysts with SQL skills
- Getting results incurs development effort!

# Plan for today

- Beyond MapReduce 
- Higher-level languages for Hadoop
  - Hive Query Language 
  - Pig and Pig Latin
- Abstractions for iterative batch-processing
  - Pregel: Bulk Synchronous Parallel for Graphs

# Hive: SQL on top of Hadoop

- SQL is a higher-level language than MapReduce
  - Problem: Company may have lots of people with SQL skills, but few with Java/MapReduce skills
- Can we “bridge the gap” somehow?

```
SELECT a.campaign_id, count(*), count(DISTINCT b.user_id)
FROM dim_ads a JOIN impression_logs b ON(b.ad_id=a.ad_id)
WHERE b.dateid = '2008-12-01'
GROUP BY a.campaign_id
```

- **Idea:** SQL frontend for MapReduce
  - Abstract delimited files as tables (give them schemas)
  - Compile (approximately) SQL to MapReduce jobs!

# Recall: Database Mgmt System

- An abstract storage system
  - Provides access to **tables**, organized however the database administrator and the system have chosen
- Relational data model
  - Schema formally describes fields, data types, and constraints
- A **declarative** processing model
  - Query language: SQL or similar
  - We describe what we want to store or compute, not how it should be done
  - More general than (single-pass) MapReduce
- A strong **consistency and durability** model
  - Transactions with ACID properties




# Roles of a DBMS

- Online transaction processing (OLTP)
  - Workload: Mostly updates
  - Examples: Order processing, flight reservations, banking, ...
- Online analytic processing (OLAP)
  - Workload: Mostly queries
  - Aggregates data on different axes; often step towards mining
- May well have combinations of both

# The database approach

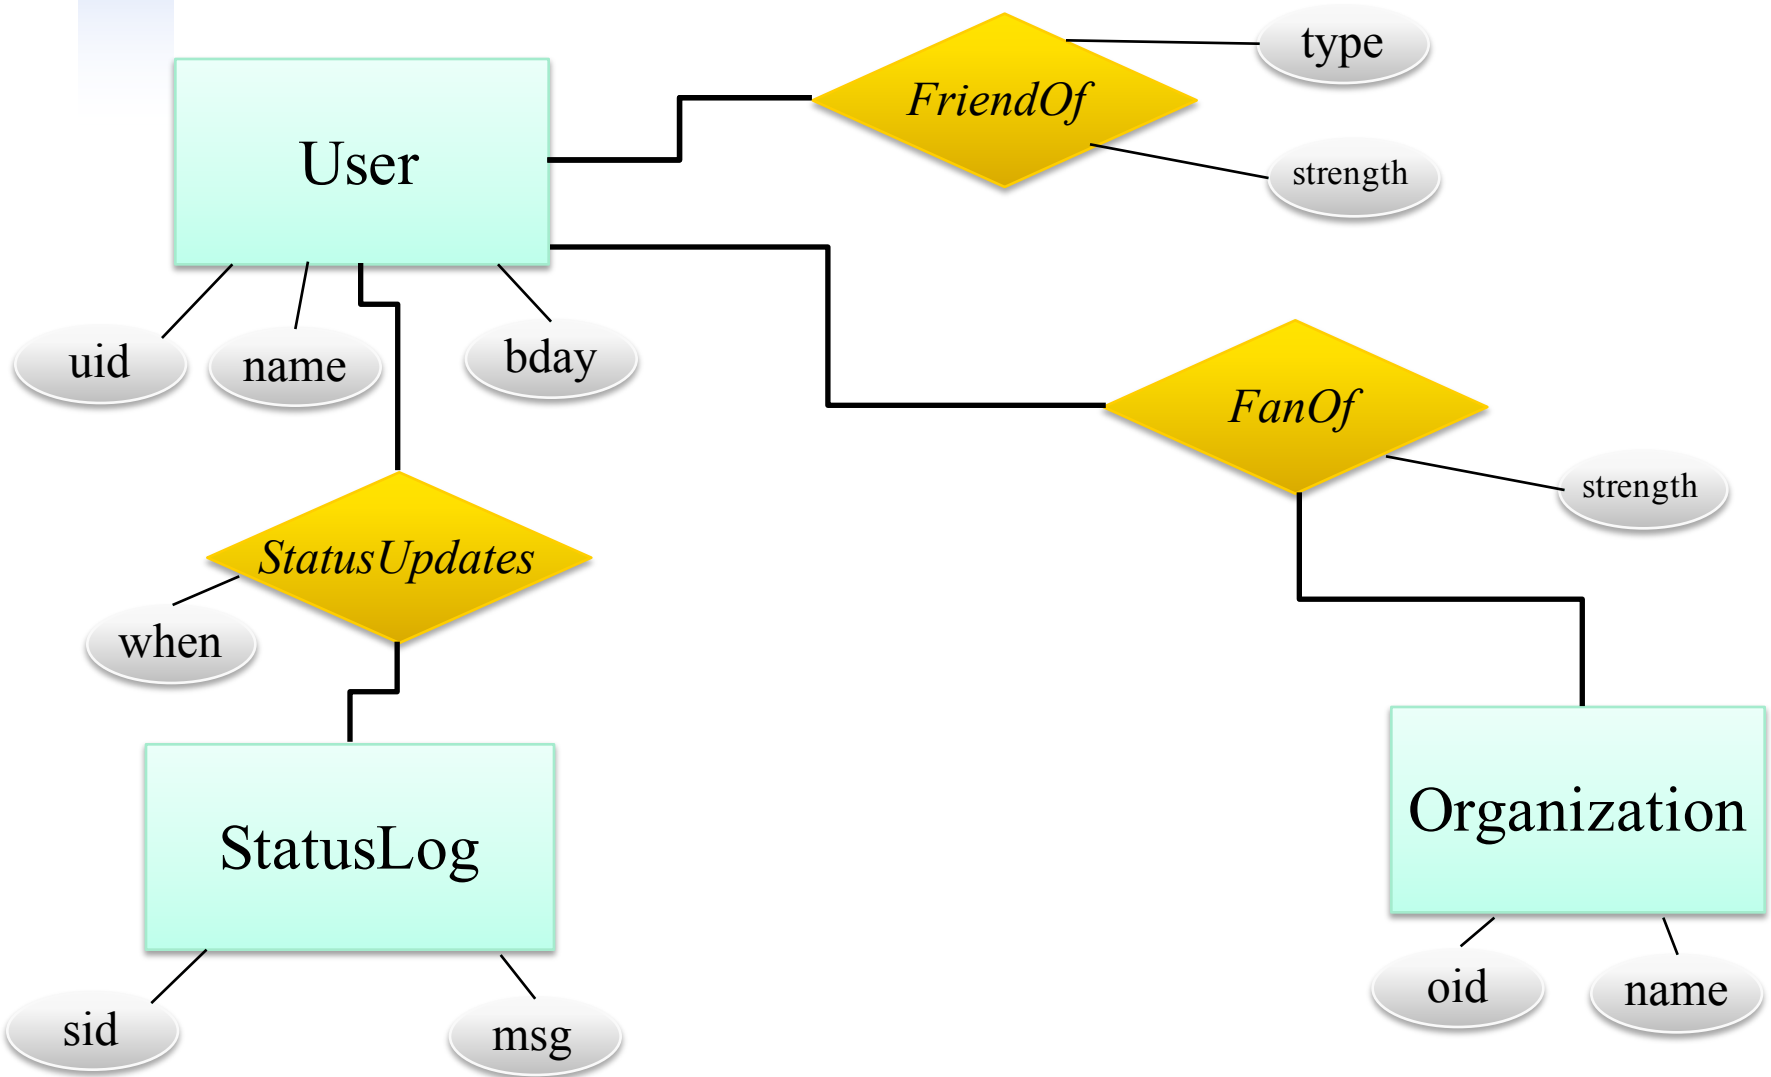
- **Idea:** User should work at a level close to the specification – not the implementation
  - A **logical** model of the data – a **schema**
    - Basically like class definitions, but also includes relationships, constraints
  - This will help us form a **physical** representation, i.e., a set of tables
    - Applications stay the same even if the platform changes
- **Computations are specified as queries**
  - Again, in terms of logical operations
  - Gets mapped into a **query evaluation plan**

# Recall: Our (simplistic) social network



(Alice, **fan-of**, 0.5, Facebook)  
(Alice, **friend-of**, 0.9, Sunita)  
(Jose, **fan-of**, 0.5, Magna Carta)  
(Jose, **friend-of**, 0.3, Sunita)  
(Mikhail, **fan-of**, 0.8, Facebook)  
(Mikhail, **fan-of**, 0.7, Magna Carta)  
(Sunita, **fan-of**, 0.7, Facebook)  
(Sunita, **friend-of**, 0.9, Alice)  
(Sunita, **friend-of**, 0.3, Jose)

# Logical schema with entity-relationship



# Some example tables

User

uid	name	bdate
1	alice	1-1-80
2	jose	1-1-70
3	sunita	6-1-75

FriendOf

uid	fid	strength	type
1	3	0.9	fr
2	3	0.3	fr

StatusUpdates

uid	sid	when
1	1	10/1
2	17	11/1

FanOf

uid	oid	strength
1	99	0.5
2	100	0.5
3	99	0.7

StatusLog

sid	post
1	In Rome
17	Drank a latte

Organization

oid	name
99	Facebook
100	Magna Carta

# Recap: Databases

- A more abstract view of the data
  - Schema formally describes fields, data types, and constraints
  - Relational model: Data is stored in tables
  - Declarative: We describe what we want to store or compute, not how it should be done
  - The implementation (a database management system, or DBMS) takes care of the details
- Much higher-level than MapReduce
  - This has both pros and cons

# Basics of querying in SQL

- At its core, a database query language consists of manipulations of sets of tuples
  - We **bind** variables to the tuples within a table, perform tests on each value, and then construct an output set
  - Java:

```
ArrayList<String> output ...  
for (u : Table<User>) { output.add(u.name); }
```
  - Map/Reduce:

```
public void map(LongWritable k, User v)  
    { context.write(new Text(v.name)); }
```
  - SQL:

```
SELECT U.name FROM User U
```

# The SQL standard form

- Each **block** computes a set/bag of tuples
- A block looks like this:

```
SELECT [DISTINCT] {T1.attrib, ..., T2.attrib}  
FROM {relation} T1, {relation} T2, ...  
WHERE {predicates}
```

```
GROUP BY {T1.attrib, ..., T2.attrib}  
HAVING {predicates}  
ORDER BY {T1.attrib, ..., T2.attrib}
```



# Multiple table variables in SQL

- Recall from a couple of slides back:

```
SELECT U.name  
FROM User U
```

returns (name) tuples

- We can compute all combinations of possible values (**Cartesian product** of tuples) as:

```
SELECT U.name, U2.name  
FROM User U, User U2
```

- Or we can compute a **union** of tuples as:

```
(SELECT U.name FROM User U)  
UNION  
(SELECT O.name FROM Organization O)
```

# The basic operations

- So far, we've seen how to combine tables
- Let's see some more sophisticated operations:
  - Filtering
  - Remapping / renaming / reorganizing
  - Intersecting
  - Sorting
  - Aggregating

# Filtering and remapping

- **Filtering** is very easy – simply add a test in the WHERE clause:

```
SELECT *  
FROM User  
WHERE name LIKE 'j%'
```

(Note \*, LIKE, %)

- We can also **reorder**, **rename**, and **project**:

```
SELECT name, uid AS id  
FROM User  
WHERE name LIKE 's%'
```

# Intersection and join

- True **intersection** – “same kind” of tuples:

```
(SELECT U.name FROM User U)
INTERSECT
(SELECT O.name FROM Organization O)
```

- **Join** – merge tuples from different table variables when they satisfy a condition:

```
SELECT U.name, S.post
FROM User U, StatusUpdates P, StatusLog S
WHERE U.uid = P.uid AND P.sid = S.sid
```

- If the attribute names are the same:

```
SELECT U.name, S.post
FROM User U NATURAL JOIN StatusUpdates SU
NATURAL JOIN StatusLog S
```

# Sorting

- Output order is arbitrary in SQL
- Unless you specifically ask for it:

```
SELECT *  
FROM USER U  
ORDER BY name
```

```
SELECT *  
FROM USER U  
ORDER BY name DESC
```

# Aggregating on a key: Group By

- What if we wanted to compute the average friendship strength per organization?

- Need to group the tuples in FanOf by 'oid', then average

- This can be done with Group By:

```
SELECT {group-attribs}, {aggregate-op} (attrib)
FROM {relation} T1, {relation} T2, ...
WHERE {predicates}
GROUP BY {group-list}
```

- Built-in aggregation operators:

- AVG, COUNT, SUM, MAX, MIN
  - DISTINCT keyword for AVG, COUNT, SUM

# Example: Group By

- Recall the k-means algorithm

- Suppose we want to compute the new centroids for a set of points, and we already have the points as a table  
PointGroups(PointID, GroupID, X, Y)

```
SELECT P.GroupID, AVG(P.X), AVG(P.Y)
FROM PointGroups P
GROUP BY P.GroupID
```

- Can also write aggregation, e.g., in C, Java

- Example: Oracle's Java Stored Procedures
- Basically like the Reduce function!

# Composition

- The results of SQL are tables
  - Hence you can query the results of a query!

- Let's do k-means in SQL:

```
SELECT PG.GroupID, AVG(PG.X), AVG(PG.Y)
FROM (
  SELECT P.ID, P.X, P.Y,
         ARGMIN(dist(P.X, P.Y, G.X, G.Y), G.ID),
         MIN(dist(P.X, P.Y, G.X, G.Y))
  FROM POINTS P, GROUPS G
  GROUP BY P.ID
) AS PG
GROUP BY PG.GroupID
```



# Recap: Querying with SQL

- We have seen SQL constructs for:
  - Projection and remapping/renaming (SELECT)
  - Cartesian product (FROM x, y, z, ...; NATURAL JOIN)
  - Filtering (WHERE)
  - Set operations (UNION, INTERSECT)
  - Aggregation (GROUP BY + MIN, MAX, AVG, ...)
  - Sorting (ORDER BY)
  - Composition (SELECT ... FROM (SELECT ... FROM ...))
- Not a complete list - SQL has more features!

# Hive

- A data warehouse infrastructure built on top of Hadoop for providing data summarization, query and analysis
- Hive Query Language (HQL) – similar to SQL
  - Suitable for processing structured data
  - Create a table structure on top of HDFS
  - Queries are compiled in to MapReduce jobs
- Not designed for OLTP!
  - Updating records or transactions are not supported

# Example: WordCount

```
CREATE TABLE doc (line STRING);  
LOAD DATA LOCAL INPATH 'text.txt' INTO TABLE doc;  
  
CREATE TABLE wordcount AS  
SELECT word, count(1) AS count  
FROM (SELECT EXPLODE(SPLIT(line, '\s')) AS word FROM doc) words  
GROUP BY word  
ORDER BY count DESC, word ASC;
```

# Plan for today

- Beyond MapReduce ✓
- Higher-level languages for Hadoop
  - Hive Query Language ✓
  - Pig and Pig Latin
- Abstractions for iterative batch-processing
  - Pregel: Bulk Synchronous Parallel for Graphs



# Towards Pig #1: Beyond relations?

- The relational data model allows us to have arbitrary numbers of **relations**
  - Each with its own schema that includes arbitrary numbers of attributes
- But: No nested tables!
  - These would be converted into multiple tables by 1NF normalization
  - Hence SQL has no nested collections at all, (sets, lists, bags...)
- Can we add support for these?

# Towards Pig #2: Programming model

## ■ Hadoop MapReduce:

- file-oriented, **procedural**
- regularized “pipeline” – map, combine, shuffle, reduce
- arbitrary Java functions at each step

rigid dataflow

opacity

custom code  
even for very

common operations

## ■ SQL:

- random access-storage-oriented (DBMS controls storage)
- compositional, tuple-collection-oriented query model
- **declarative** queries are automatically optimized
- can accommodate Java functions, but not naturally
- Hive: SQL queries → file-oriented Map/Reduce

what about  
"procedural  
programmers"?

## ■ Is there something in between?

- Declarative is nice, but many data analysts are 'entrenched' procedural programmers...
- Pig and Pig Latin!

# Pig Latin and Pig

- **Pig Latin**: a compositional, collections-oriented **dataflow** language
  - Oriented towards parallel data processing & analysis
  - Think of it as a more procedural SQL-like language with nested collections
    - Emphasizes **user-defined functions**, esp. those that have nice algebraic properties (unlike SQL)
    - Supports external data from files (like Hive)
  - By Chris Olston et al. at Yahoo! Research
    - [http://www.tomkinshome.com/site\\_media/papers/papers/ORS+08.pdf](http://www.tomkinshome.com/site_media/papers/papers/ORS+08.pdf)
- **Pig**: the runtime system

# Pig Latin: Basic constructs

- Collection-valued expressions whose results get assigned to variables
  - A program does a series of assignments in a dataflow
  - It gets compiled down to a sequence of MapReduces
    - Similar to Hive, but Pig Latin has its own query language (not SQL)
- Basic SQL-like operations are explicitly specified:
  - `load ... as` [HDFS scan]
  - **Remapping:** `foreach ... generate` [Map]
  - **Filtering:** `filter by` [Map]
  - **Intersecting:** `join` [Reduce]
  - **Aggregating:** `group by` [Reduce]
  - **Sorting:** `order` [Shuffle]
  - `store` [HDFS store]

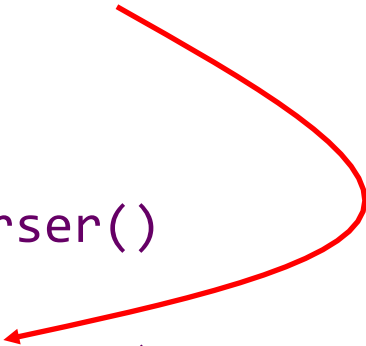


# Simple example: Face detection

- Each expression creates a **named collection**
  - load collections from files
  - process them (e.g., per tuple) using a user-defined function
  - store the results into files

```
I = load '/mydata/images' using ImageParser()  
    as (id, image);
```

```
F = foreach I generate id, detectFaces(image);  
store F into '/mydata/faces';
```



# Example: Session Classification

- Goal: Find web sessions that end on the 'best' page (i.e., the page with the highest PageRank)
  - We need to join two tables, and then compare the final rank in the sequence to the other ranks

Visits

User	URL	Time
Alice	<a href="http://www.cnn.com">www.cnn.com</a>	7:00
Alice	<a href="http://www.digg.com">www.digg.com</a>	7:20
Alice	<a href="http://www.social.com">www.social.com</a>	10:00
Alice	<a href="http://www.flickr.com">www.flickr.com</a>	10:05
Joe	<a href="http://www.cnn.com/index.htm">www.cnn.com/index.htm</a>	12:00

⋮

Pages

URL	PageRank
<a href="http://www.cnn.com">www.cnn.com</a>	0.9
<a href="http://www.flickr.com">www.flickr.com</a>	0.9
<a href="http://www.social.com">www.social.com</a>	0.7
<a href="http://www.digg.com">www.digg.com</a>	0.2

⋮

# The computation in Pig Latin

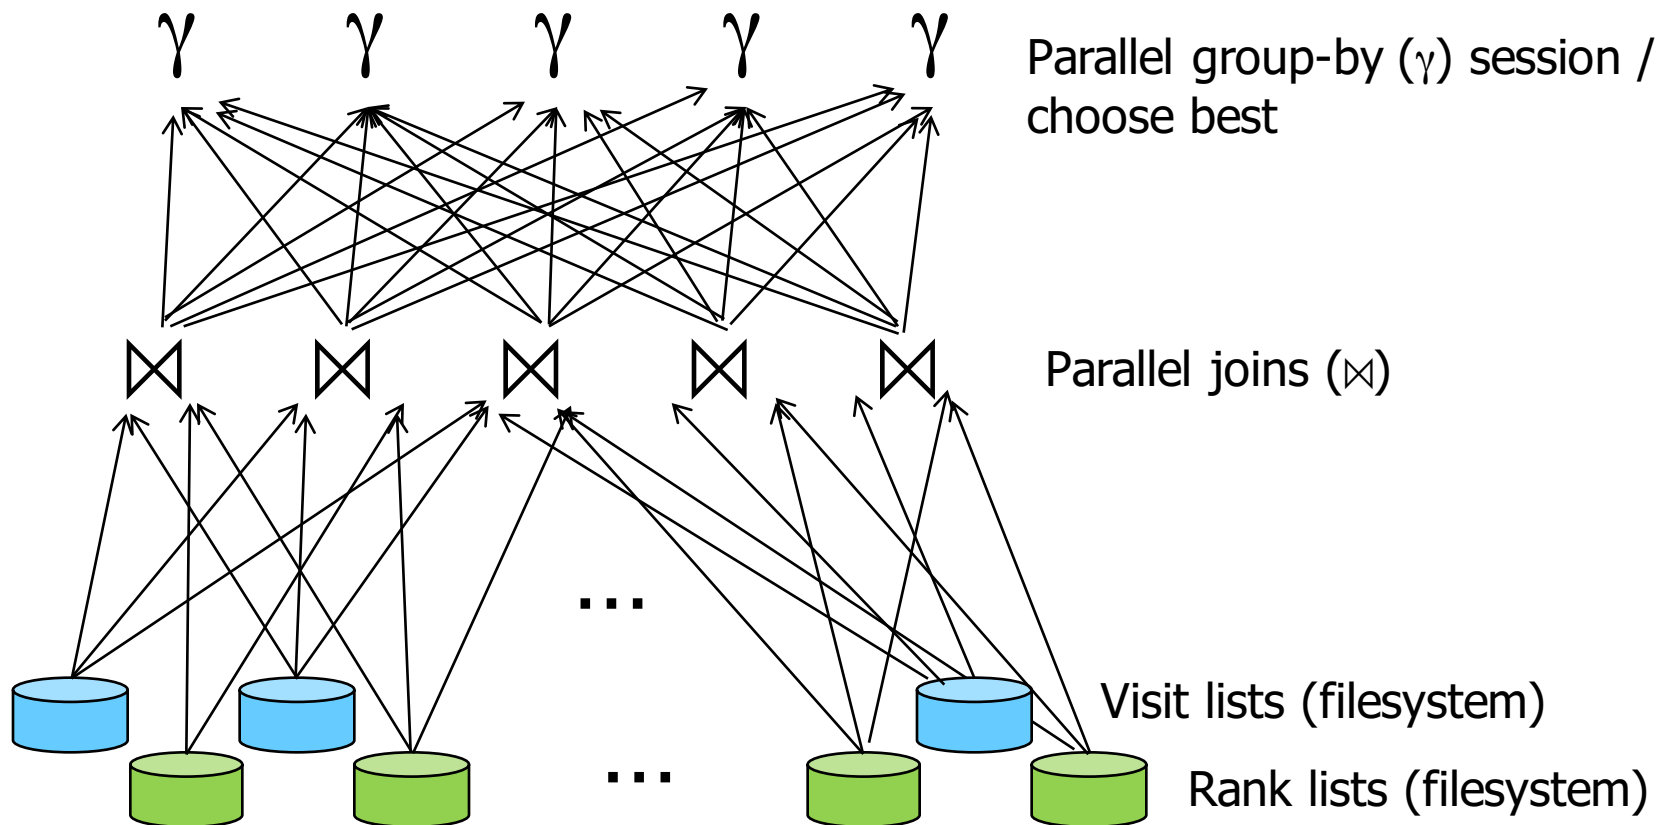
```
Visits = load '/data/visits' as (user, url, time);  
Visits = foreach Visits generate user, Canonicalize(url),  
    time;
```

```
Pages = load '/data/pages' as (url, pagerank);
```

```
    VP = join Visits by url, Pages by url;  
UserVisits = group VP by user;  
    Sessions = foreach UserVisits generate  
        flatten(FindSessions(*));  
HappyEndings = filter Sessions by BestIsLast(*);  
  
    store HappyEndings into '/data/happy_endings';
```

# What does this query compile to?

- Parallel evaluation is really a Map-Map/Reduce/Reduce chain:



# Pig Latin features

- Record-oriented transformations
  - Can work over, create nested collections
  - (Resembles Nested Relational variants of SQL)
- Basic operators expose parallelism; user-defined operators may not
- Operations are explicit, not declarative
  - Unlike SQL

operators:

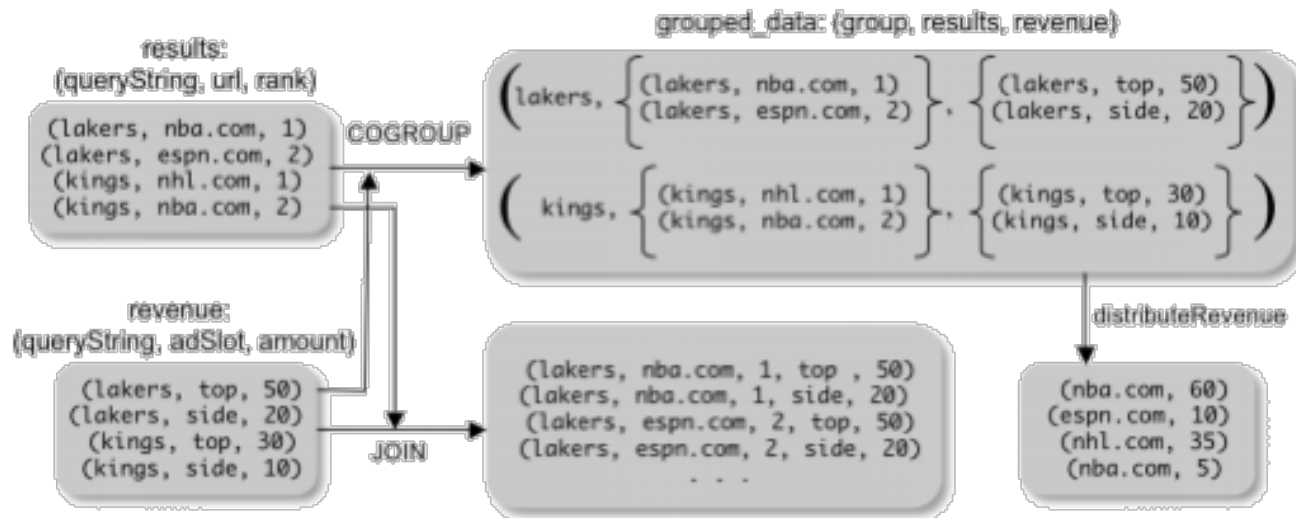
- FILTER
- FOREACH ... GENERATE
- GROUP

binary operators:


- JOIN
- COGROUP
- UNION

# Nesting: COGROUP & FLATTEN

- Cogrouping: nesting groups into columns



- Flattening: unnesting groups



# Pig Latin vs. MapReduce

- MapReduce combines 3 primitives:  
process records → create groups → process groups

```
a = FOREACH input GENERATE flatten(Map(*));  
b = GROUP a BY $0;  
c = FOREACH b GENERATE Reduce(*);
```

- In Pig, these primitives are:
  - explicit
  - independent
  - fully composable
- Pig adds primitives for:
  - filtering tables
  - projecting tables
  - combining 2 or more tables

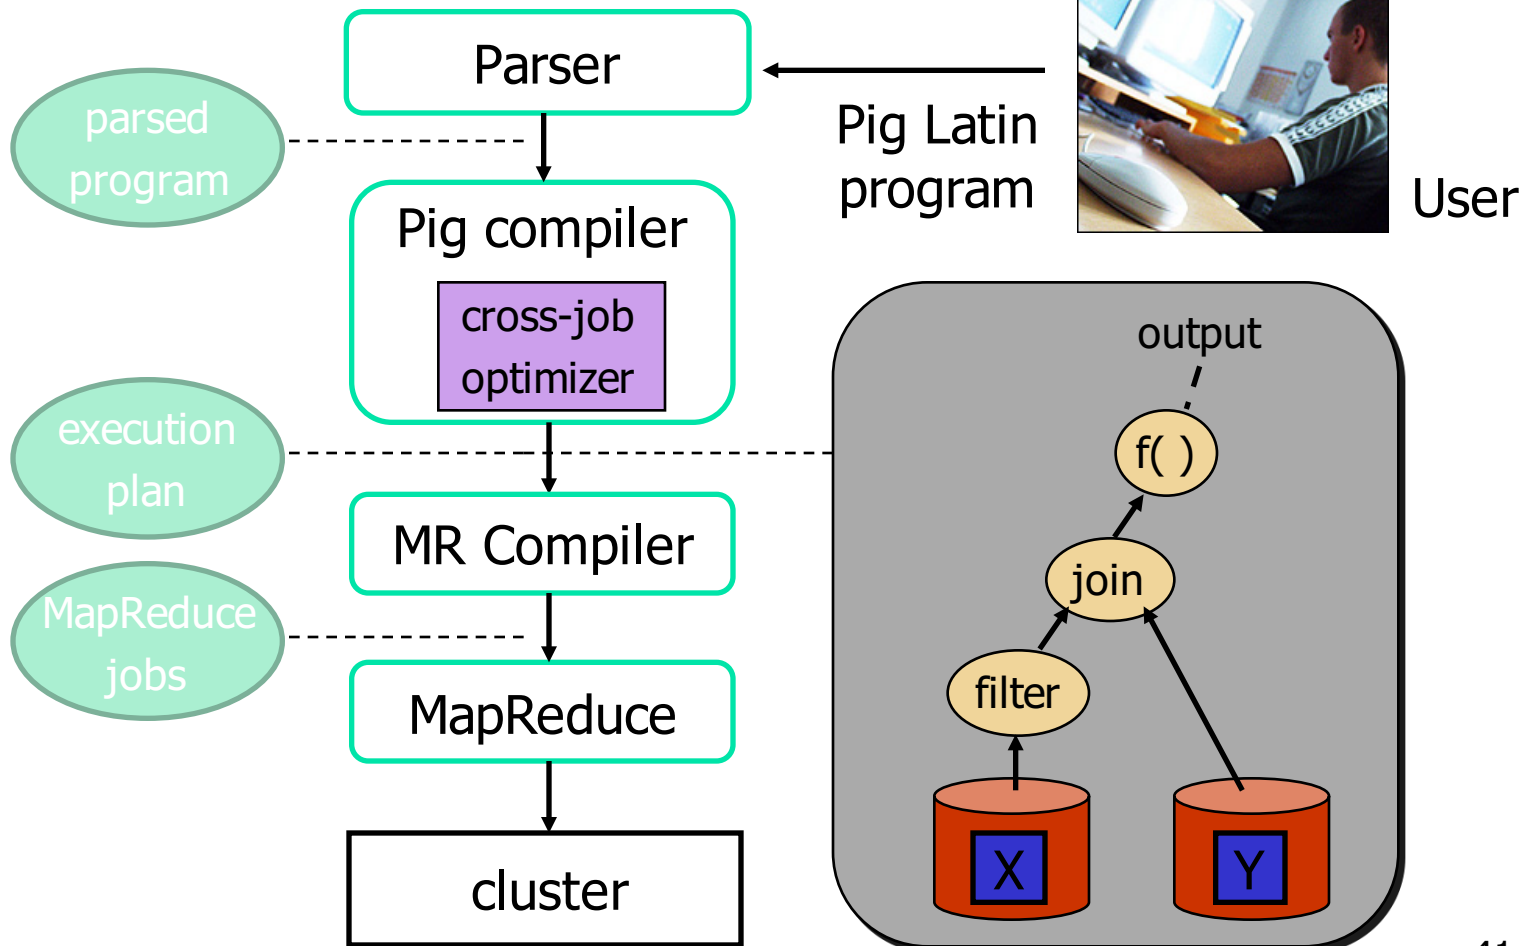
# Recap: Pig Latin

- A dataflow language that compiles to MapReduce
  - Borrows many of the elements of SQL, but eliminates the reliance on declarative optimization
  - Incorporates primitives for nested collections
- Quite successful:
  - As of 2008: 25% of Yahoo Map/Reduce jobs from Pig
  - Part of the Hadoop standard distribution



# Pig system implementation

- Let's briefly look at the Pig implementation, and how it can do a bit more because of the higher-level language:




# Key issue: Minimizing redundancy


- Popular tables
  - web crawl
  - search log
- Popular transformations
  - eliminate spam pages
  - group pages by host
  - join web crawl with search log
- Goal: Minimize redundant work

# Work-sharing techniques


execute similar jobs together



cache data transformations



cache data moves



# Recap: Pig and Pig Latin

- Somewhere between a programming language and a DBMS
- Allows distributed programming with explicit parallel dataflow operators
- Supports explicit management of nested collections
- Runtime system does caching and batching

# Plan for today

- Beyond MapReduce ✓
- Higher-level languages for Hadoop
  - Hive Query Language ✓
  - Pig and Pig Latin ✓
- Abstractions for iterative batch-processing
- Pregel: Bulk Synchronous Parallel for Graphs



# New Abstractions Needed

Much of the mismatch stems from the lack of shared global state

Complex applications and interactive queries both need one thing that MapReduce lacks

- Efficient primitives for data sharing

# What If We Could Remember?

Suppose we were to change things entirely:

- A set of machines
- ... each with a *partition* of a dataset, stored in memory
- Computation consists of *sending updates* from one portion to another


Let's look at two versions of this

# Pregel: Bulk Synchronous Parallel

Let's slightly rethink the MapReduce model for processing **graphs**

- Vertices
- "Edges" are really messages

Compare to MapReduce keys  $\rightarrow$  values?



"Think like a vertex"




# The Basic Pregel Execution Model

A sequence of *supersteps*, for each vertex  $V$


At superstep  $S$ :

- Compute in parallel at each  $V$ 
  - Read messages sent to  $V$  in superstep  $S-1$
  - Update value / state
  - Optionally change topology
- Send messages
- Synchronization
  - Wait till all communication is finished




# Termination Test

- Based on every vertex voting to halt
  - Once a vertex deactivates itself it does no further work unless triggered externally by receiving a message
- Algorithm terminates when all vertices are simultaneously inactive

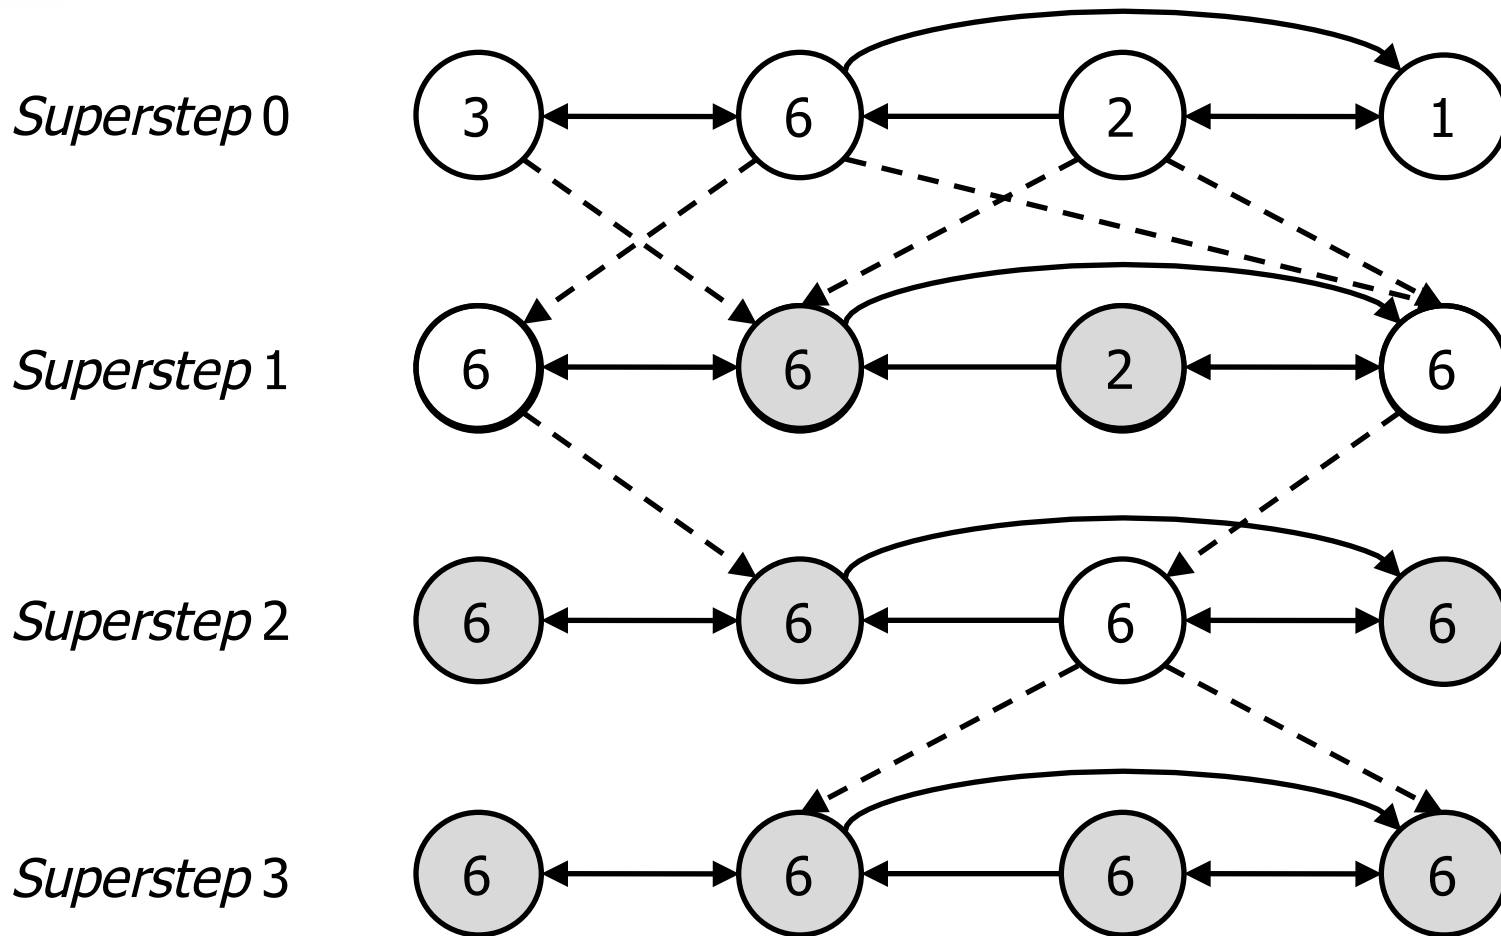


# Example: Find Maximum Value



- Each vertex contains an integer value
- Idea: propagate the largest value to every vertex
  - At superstep 0, start by propagating value to all neighbors
  - In each superstep, any vertex that has learned a larger value from its messages sends it to all its neighbors; otherwise vote to halt
  - Terminates when no further vertices change in a superstep

# Example: Find Maximum Value



# Recall: PageRank

- Each page  $i$  is given a rank  $x_i$
- Goal: Assign the  $x_i$  such that the rank of each page is governed by the ranks of the pages linking to it:

$$x_i = \sum_{j \in B_i} \frac{1}{N_j} x_j$$

Rank of page  $i$

Rank of page  $j$

Number of links out from page  $j$

Every page  $j$  that links to  $i$

How do we compute the rank values?

# Naïve PageRank Algorithm Restated

- Let

- $N(p)$  = number outgoing links from page  $p$
- $B(p)$  = number of back-links to page  $p$

$$PageRank(p) = \sum_{b \in B(p)} \frac{1}{N(b)} PageRank(b)$$

- Each page  $b$  distributes its importance to all of the pages it points to (so we scale by  $1/N(b)$ )
- Page  $p$ 's importance is increased by the importance of its back set
- Iterate till convergence or some number of iterations

# PageRank in Pregel

```
void Compute(messages) {  
    if (superstep() >= 1) {  
        sum = 0;  
        foreach (msg in messages)  
            sum += msg->Value();  
        value = 0.15 / NumVertices() + 0.85 * sum;  
        SetValue(value);  
    }  
    if (superstep() < 30) {  
        n = GetNumOfOutEdges();  
        SendMessageToAllNeighbors(GetValue() / n);  
    } else {  
        VoteToHalt();  
    }  
}
```

# Pregel Summary

- Bulk Synchronous Parallel – sequence of synchronized supersteps
  - Abstraction originally invented by Leslie Valliant in the '80s
- Consider the nodes to have state (memory) that carries from superstep to superstep
- Connections to MapReduce model?
- See also Apache Hama, Giraph, Graph.lab



# Stay tuned



Next time you will learn about:

**Beyond MapReduce – In-memory processing, Streaming**