

# INGI2145: CLOUD COMPUTING (Fall 2015)

Beyond MapReduce: In-memory processing, Streaming

5 November 2015

# MapReduce: Not for Every Task

- MapReduce greatly simplified large-scale data analysis on unreliable clusters of computers
  - Brought together many traditional CS principles
    - functional primitives; master/slave; replication for fault tolerance
  - Hadoop adopted by many companies
  - Affordable large-scale batch processing for the masses

But increasingly people wanted more!

# MapReduce: Not for Every Task

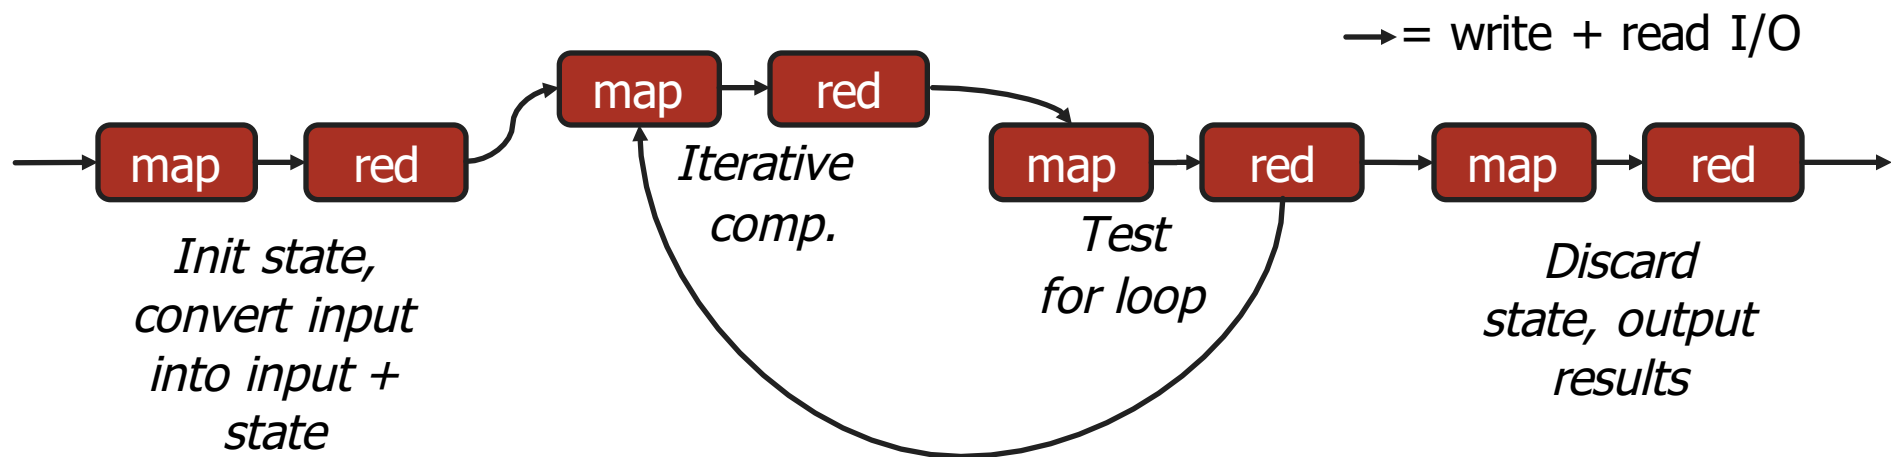
But increasingly people wanted more:

- More complex, multi-stage applications
- More interactive ad-hoc queries
- Process live data at high throughput and low latency

Which are not a good fit for MapReduce...

# MapReduce for Iterative Computation

- MapReduce is essentially functional
- Expressing iterative algorithms as chains of Map/Reduce requires passing the entire state and doing a lot of network and disk I/O
  - Recall all between-stage results are materialized to reliable and distributed storage (HDFS)



# MapReduce for Ad-hoc Queries

- MapReduce specifically designed for batch operations over large amounts of data
- New analysis task means writing a new MapReduce program
  - Tedious thing to do with languages such as Java
  - Programming interface is not familiar to traditional data analysts with SQL skills
- Getting results incurs development effort!

# Plan for today

- Beyond MapReduce ✓
- Abstractions for iterative batch-processing
  - Pregel: Bulk Synchronous Parallel for Graphs ✓
  - Spark: In-Memory Resilient Distributed Datasets
- Stream processing
  - Storm: One-record at a time
  - Spark Streaming: Micro-batching



# Spark: Resilient Distributed Datasets



- Let's think of just having a big block of RAM, partitioned across machines...
  - And a series of operators that can be executed in parallel across the different partitions
- That's basically Spark
  - A distributed memory abstraction that is both fault-tolerant and efficient

# Spark: Resilient Distributed Datasets

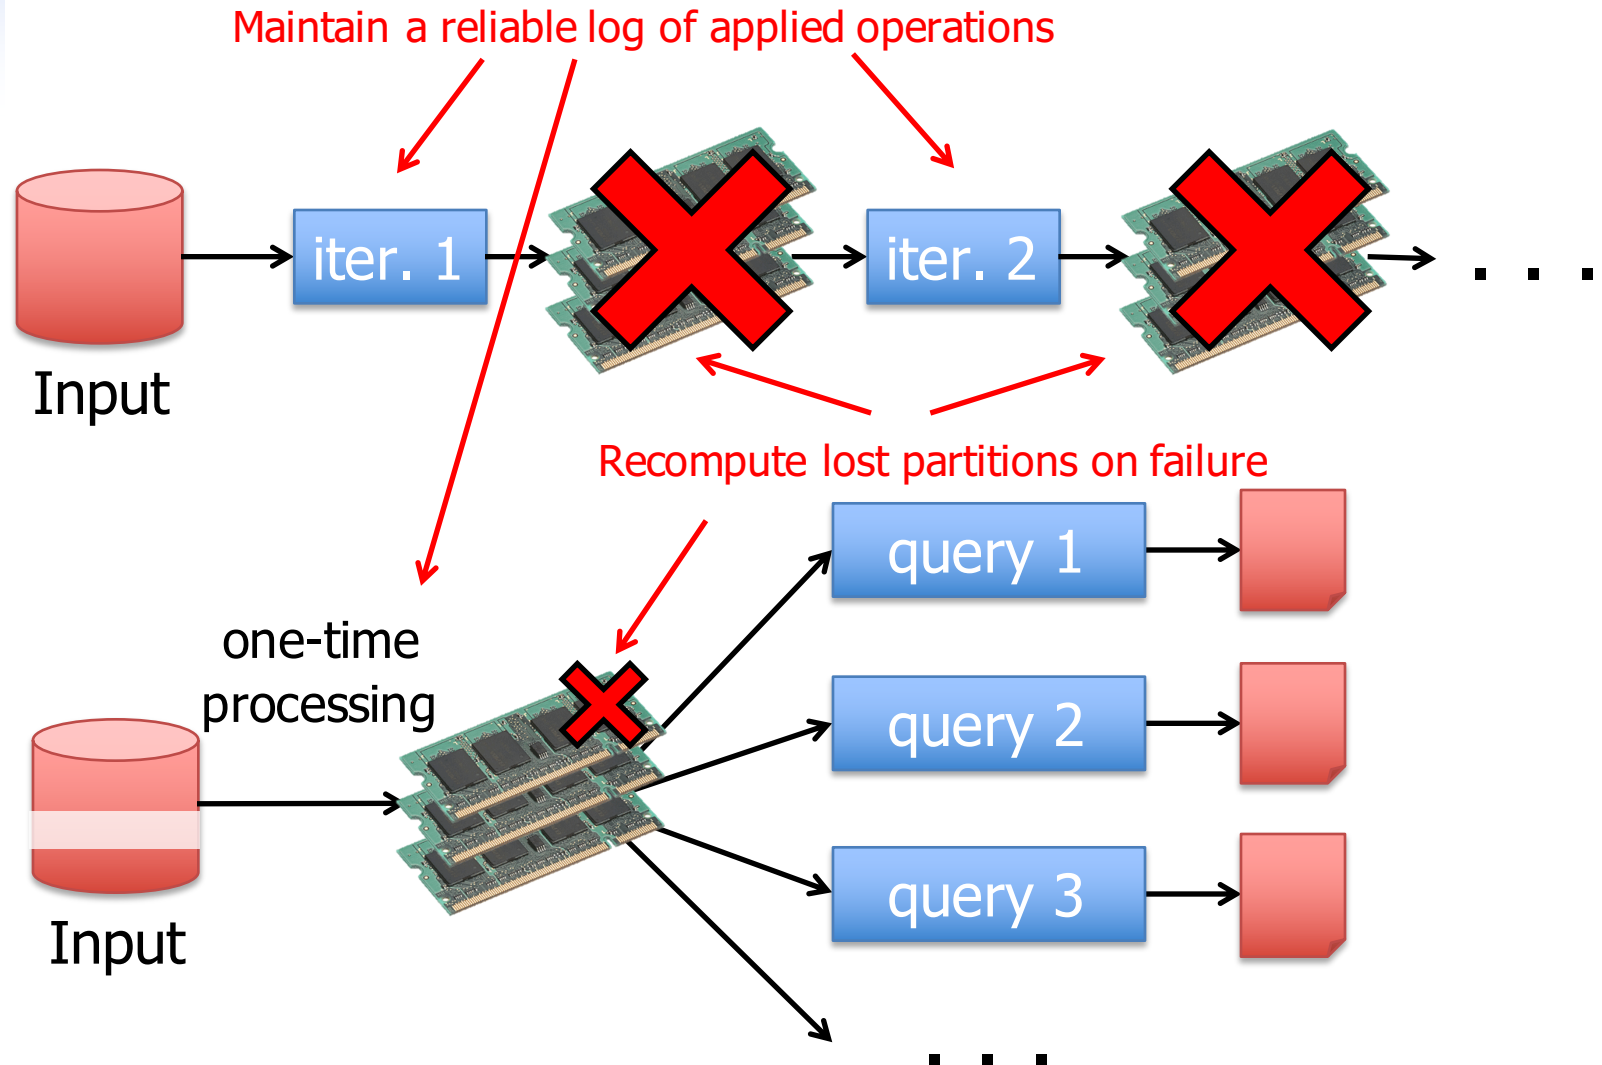
- Restricted form of distributed shared memory
  - Immutable, partitioned collections of records
  - Can only be built through *coarse-grained* deterministic transformations (map, filter, join, ...)
  - They are called Resilient Distributed Datasets (RDDs)
- Efficient fault recovery using *lineage*
  - Log one operation to apply to many elements
  - Recompute lost partitions on failure
  - No cost if nothing fails



# In-Memory Data Sharing



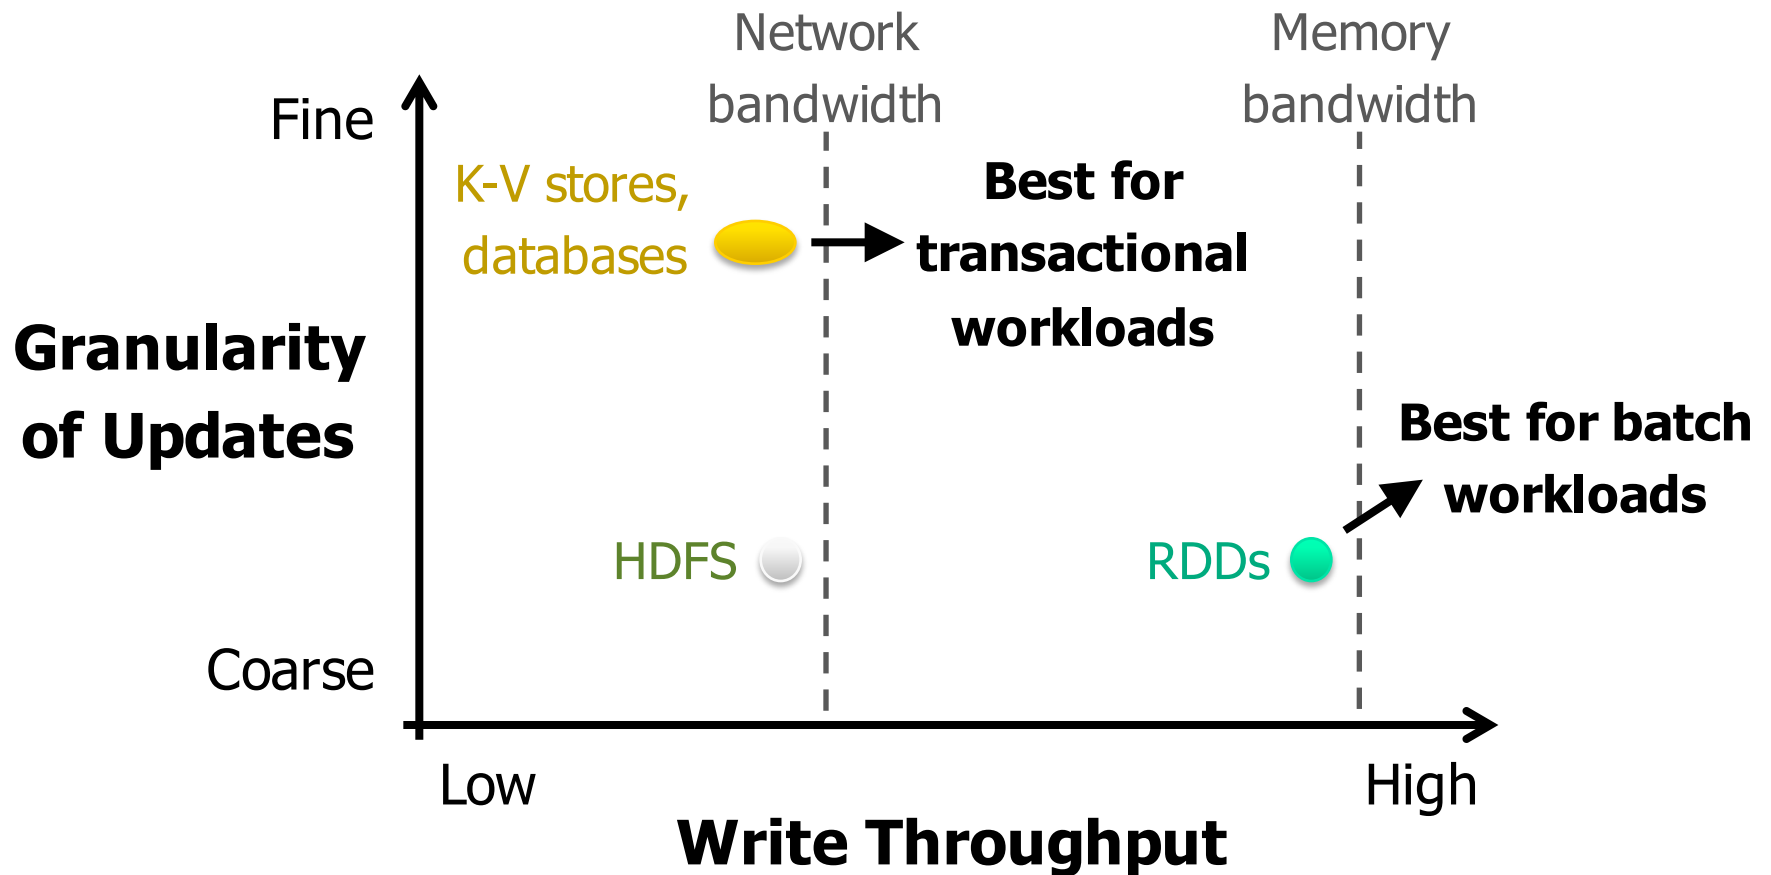
# Efficient Fault Recovery via Lineage



# Generality of RDDs

- Despite their restrictions, RDDs can express many parallel algorithms
  - These naturally *apply the same operation to many items*
- Unify many programming models
  - *Data flow models*: MapReduce, Dryad, SQL, ...
  - *Specialized models* for iterative apps: BSP (Pregel), iterative MapReduce (Haloop), bulk incremental, ...
- Support *new apps* that these models don't

# Tradeoff Space



# Spark Programming Interface

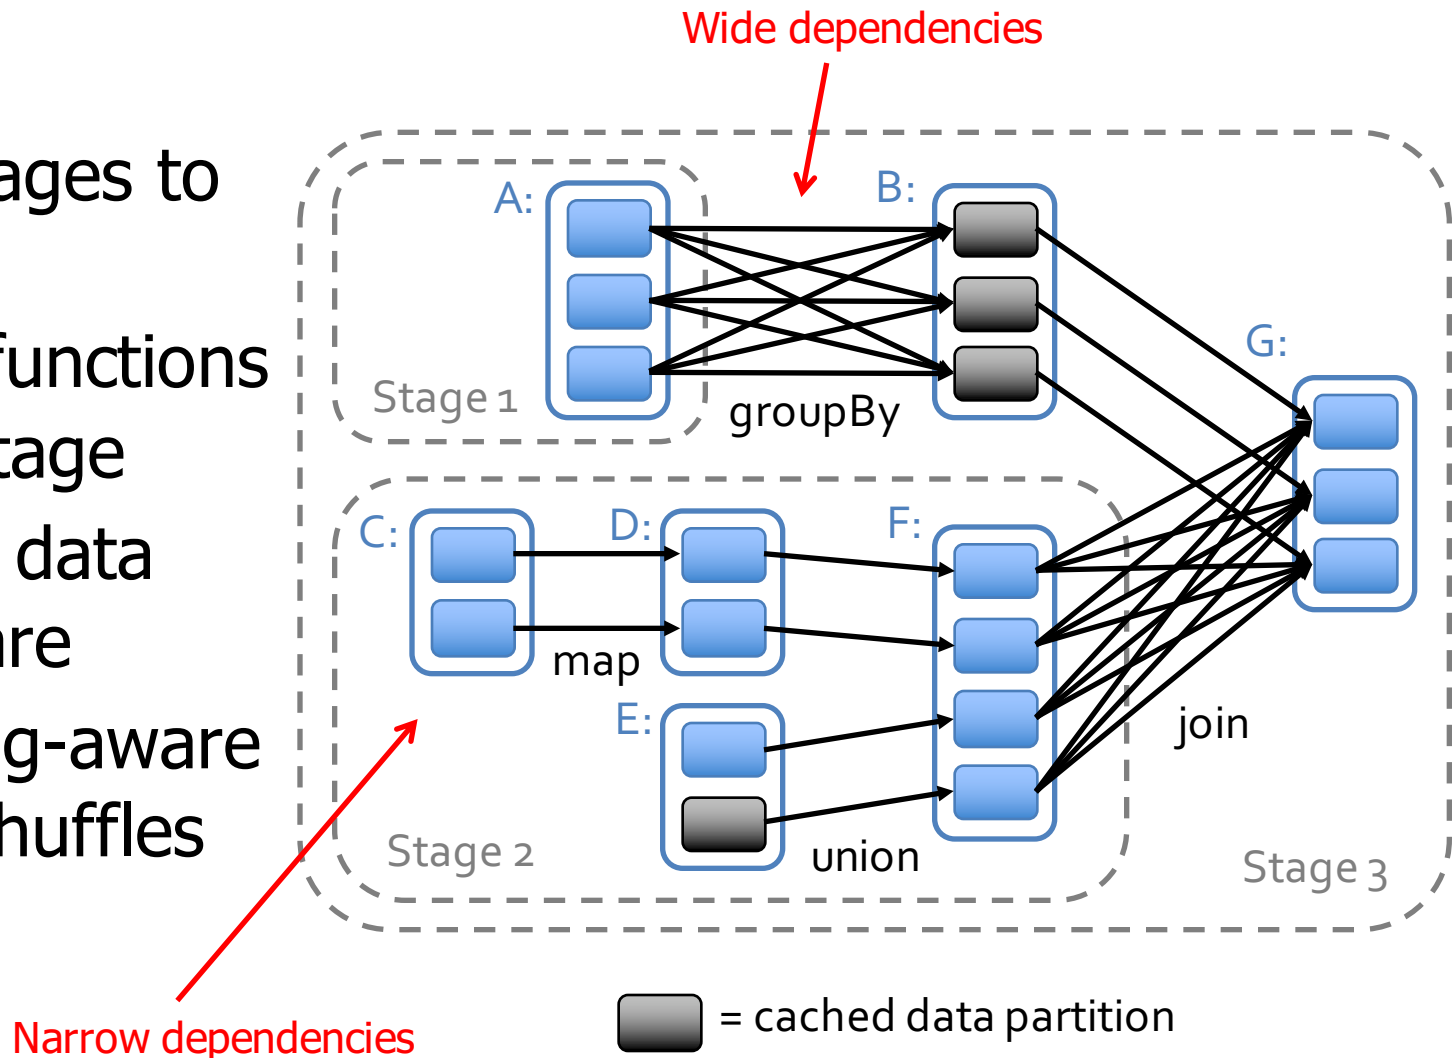
- Language-integrated API in Scala
- Usable interactively from Scala interpreter
- Provides:
  - Resilient distributed datasets (RDDs)
  - Operations on RDDs: deterministic *transformations* (build new RDDs), *actions* (compute and output results)
  - Control of each RDD's *partitioning* (layout across nodes) and *persistence* (storage in RAM, on disk, etc)

# Spark Operations

<b>Transformations</b> (define a new RDD)	map filter sample groupByKey reduceByKey sortByKey	flatMap union join cogroup cross mapValues
<b>Actions</b> (return a result to driver program)	collect reduce count save lookupKey take	

# Task Scheduler

- DAG of stages to execute
- Pipelines functions within a stage
- Locality & data reuse aware
- Partitioning-aware to avoid shuffles

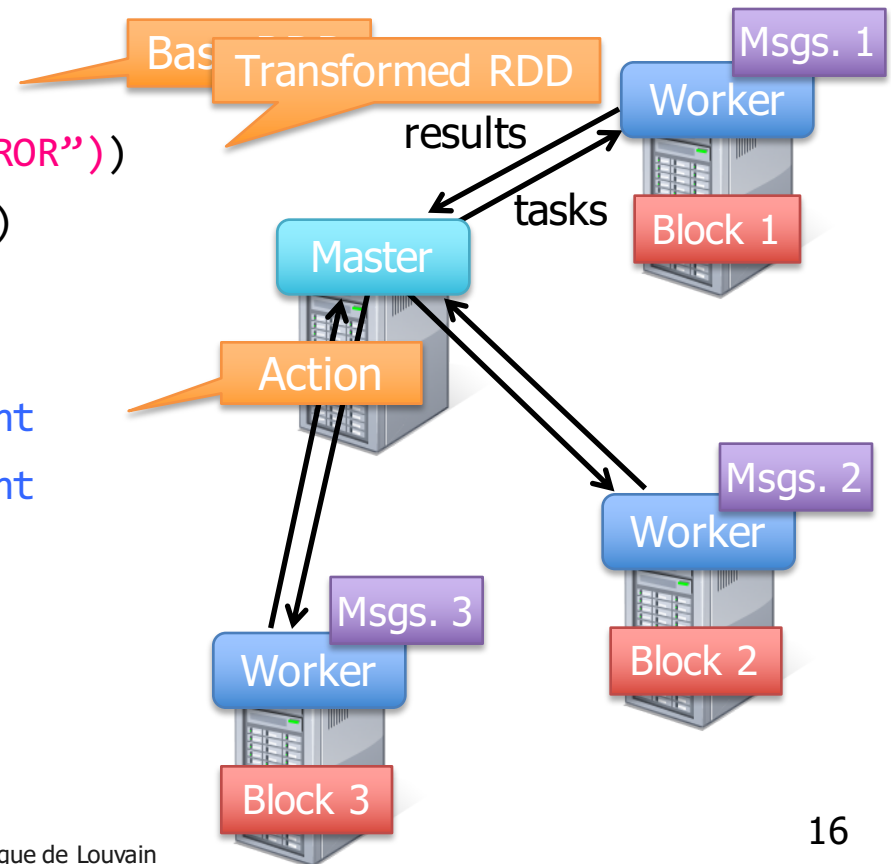


# Example: Log Mining

- Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.persist()

messages.filter(_.contains("foo")).count
messages.filter(_.contains("bar")).count
```





# Example: Word Count

## MapReduce way

```
public static class WordCountMapClass extends MapReduceBase
    implements Mapper<LongWritable, Text, Text, IntWritable> {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value,
        OutputCollector<Text, IntWritable> output,
        Reporter reporter) throws IOException {
        String line = value.toString();
        StringTokenizer itr = new StringTokenizer(line);
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            output.collect(word, one);
        }
    }
}

public static class WordCountReduce extends MapReduceBase
    implements Reducer<Text, IntWritable, Text, IntWritable> {
    public void reduce(Text key, Iterator<IntWritable> values,
        OutputCollector<Text, IntWritable> output,
        Reporter reporter) throws IOException {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

## Spark way

```
val spark = new SparkContext(master, appName,
    [sparkHome], [jars])
val file = spark.textFile("hdfs://...")
val counts = file.flatMap(line => line.split(" "))
    .map(word => (word, 1))
    .reduceByKey((x, y) => x + y)
counts.saveAsTextFile("hdfs://...")
```

# Example: Word Count

```
val spark = new SparkContext(master, appName, [sparkHome], [jars])
val file = spark.textFile("hdfs://...")
val counts = file.flatMap(line => line.split(" "))
                  .map(word => (word, 1))
                  .reduceByKey((x, y) => x + y)
counts.saveAsTextFile("hdfs://...")
```

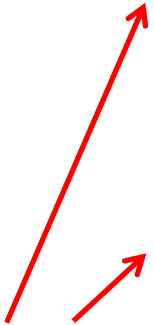
# Example: PageRank

- Iterative computation

$$PageRank(p) = (1 - d) + d \sum_{b \in B_p} \frac{1}{N(b)} PageRank(b)$$

```
var links = // RDD of (url, neighbors) pairs
var ranks = // RDD of (url, rank) pairs
```


```
for (i <- 1 to ITERATIONS) {
  contribs = links.join(ranks).flatMap {
    (url, (links, rank)) =>
      links.map(dest => (dest, rank/links.size))
  }
  ranks = contribs.reduceByKey((x, y) => x + y)
    .mapValues(sum => 0.85*sum + 0.15/N)
}
```



RDDs are immutable

contribs and ranks are new RDDs!

# Optimizing Placement



- links & ranks repeatedly joined
- Can *co-partition* them (e.g. hash both on URL) to avoid shuffles
- Can also use app knowledge, e.g., hash on DNS name
- `links = links.partitionBy(new URLPartitioner())`

# Programming Models Implemented on Spark

- RDDs can express many existing parallel models

- **MapReduce, DryadLINQ**
- **Pregel** graph processing
- **Iterative MapReduce**
- **SQL**: Hive on Spark (Shark)


All are based on  
coarse-grained  
operations

- Enables apps to efficiently *intermix* these models

# Spark Summary

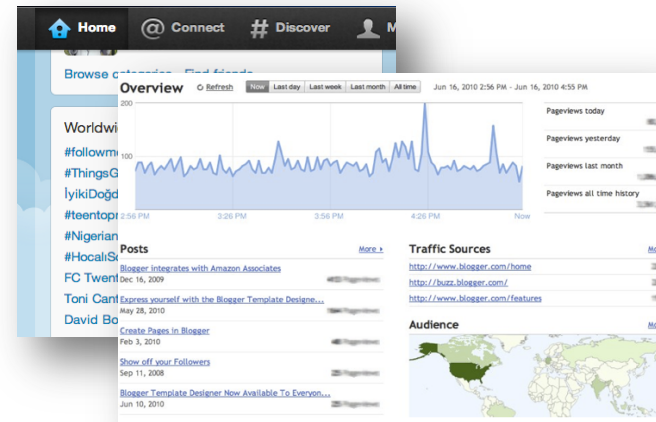
- Global aggregate computations that produce program state
  - compute the count() of an RDD, compute the max diff, etc.
- Loops!
  - Spark makes it much easier to do multi-stage MapReduce
- Built-in abstractions for some other common operations like joins
- See also Apache Crunch / Google FlumeJava for a very similar approach

# Plan for today

- Beyond MapReduce ✓
- Abstractions for iterative batch-processing ✓
  - Pregel: Bulk Synchronous Parallel for Graphs ✓
  - Spark: In-Memory Resilient Distributed Datasets ✓
- Stream processing
  - Storm: One-record at a time 
  - Spark Streaming: Micro-batching

# Stream Processing


- Many important applications must process large streams of live data and provide results in near-real-time
  - Social network trends
  - Website statistics
  - Ad impressions
- ...
- Distributed stream processing framework is required to
  - Scale to large clusters (100s of machines)
  - Achieve low latency (few seconds)






# Stateful Stream Processing

- Traditional streaming systems have a **record-at-a-time** processing model
  - Each node has mutable state
  - For each record, update state and send new records
- State is lost if node dies!
- Making stateful stream processing be fault-tolerant is challenging




# Apache Storm

- Framework for distributed stream processing
- Provides: Stream Partitioning + Fault Tolerance + Parallel Execution




Programming Model



Architecture

# Abstractions in Storm

Topology



Arbitrarily complex  
multi-stage stream  
computation

Stream



Unbounded sequence  
of tuples

Spout



Source of streams

Bolt



Process input streams and  
produce new streams  
Holds most  
computation logic

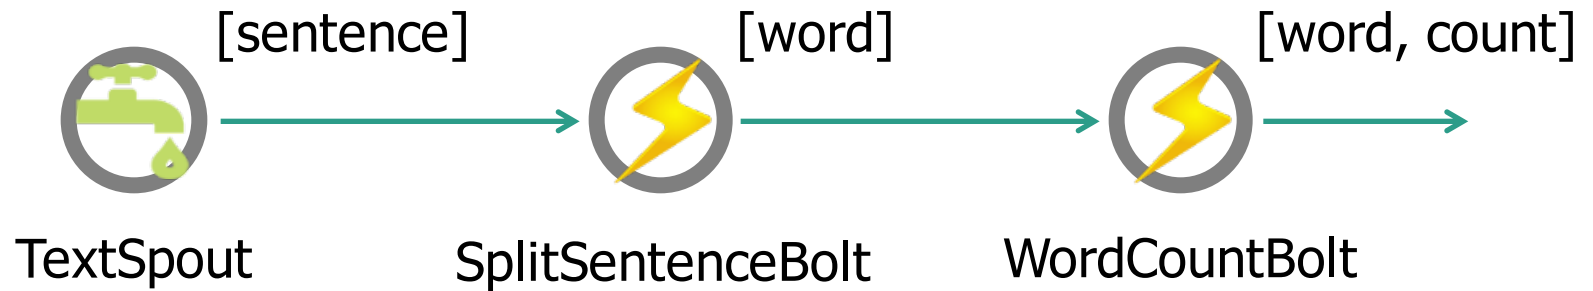
# Storm Cluster Architecture

- A storm cluster has three sets of nodes:
- Nimbus node (master node)
  - Similar to the Hadoop JobTracker
  - Distributes code, launches workers across the cluster
  - Monitors computation and reallocates workers as needed
- ZooKeeper nodes – (coordinate the cluster)
  - Will discuss ZooKeeper in detail in a later lecture
- Supervisor nodes
  - Start and stop workers according to signals from Nimbus

# Fault Tolerance


- If a supervisor node fails, Nimbus reassigns that node's task to other nodes in the cluster
- Any tuples sent to a failed node will time out and be replayed
  - Delivery guarantee dependent on a reliable data source
    - It can replay a message if processing fails at any point
- Storm can guarantee that every tuple will be process **at least once** or **at most once**, but not **exactly once**
  - Exactly once guarantee requires a durable data source that can replay any message or set of messages given the necessary selection criteria

# Example: Word Count



```
TextSpout implements IRichSpout {  
  nextTuple() {  
    while ((str = reader.readLine()) != null)  
      collector.emit(new Values(str), str);  
  }  
  [...]  
}
```

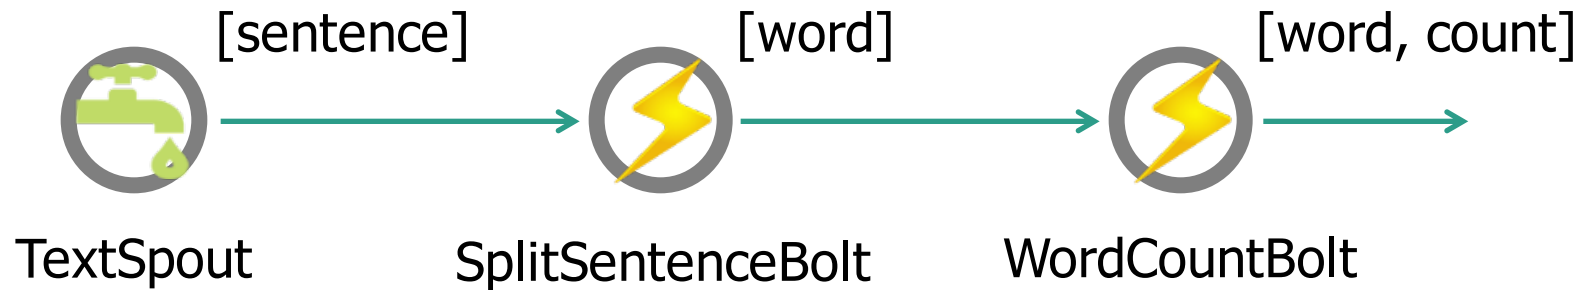
# Example: Word Count



```
class SplitSentenceBolt implements
IRichBolt {
  execute(Tuple input) {
    String sentence = input.getString(0);
    String[] words = sentence.split(" ");
    for (String word: words) {
      collector.emit(new Values(word));
    }
    collector.ack(input);
  }
  [...] }
```

```
class WordCounterBolt implements
IRichBolt {
  Map<String, Integer> counters;
  execute(Tuple input) {
    String str = input.getString(0);
    if(!counters.containsKey(str))
      counters.put(str, 1);
    else {
      Integer c = counters.get(str) + 1;
      counters.put(str, c);
    }
    collector.ack(input);
  }
  [...] }
```

# Example: Word Count



```
public class WordCountTopology {  
    [...] main(String[] args) throws Exception {  
        Config config = new Config();  
        config.setDebug(true);  
        TopologyBuilder builder = new TopologyBuilder();  
        builder.setSpout("textspout", new LineReaderSpout());  
        builder.setBolt("splitsentence", new WordSpitterBolt()).shuffleGrouping("textspout");  
        builder.setBolt("word-count", new WordCounterBolt()).shuffleGrouping("splitsentence");  
        LocalCluster cluster = new LocalCluster();  
        cluster.submitTopology("WordCountTopology", config, builder.createTopology());  
        Thread.sleep(10000);  
        cluster.shutdown();  
    } }  
}
```



# Plan for today


- Beyond MapReduce ✓
- Abstractions for iterative batch-processing ✓
  - Pregel: Bulk Synchronous Parallel for Graphs ✓
  - Spark: In-Memory Resilient Distributed Datasets ✓
- Stream processing
  - Storm: One-record at a time ✓
  - Spark Streaming: Micro-batching



# Spark Streaming

Run a streaming computation as a **series of very small, deterministic batch jobs**

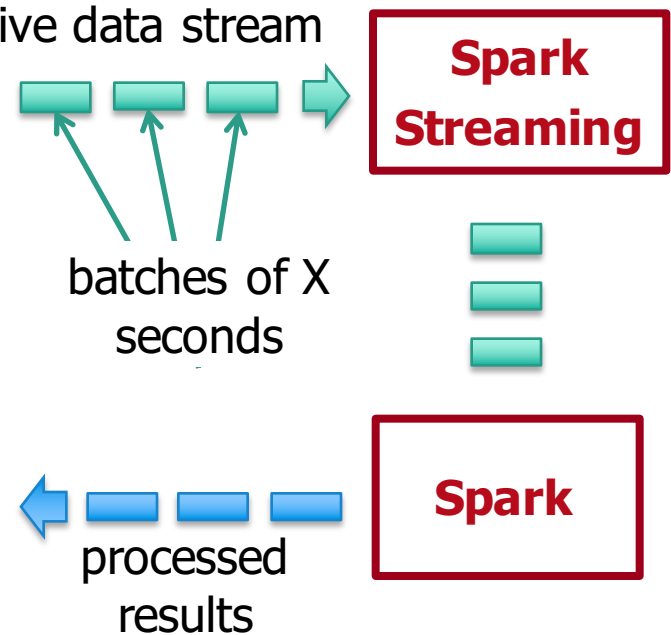
- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



# Spark Streaming

Run a streaming computation as a **series of very small, deterministic batch jobs**

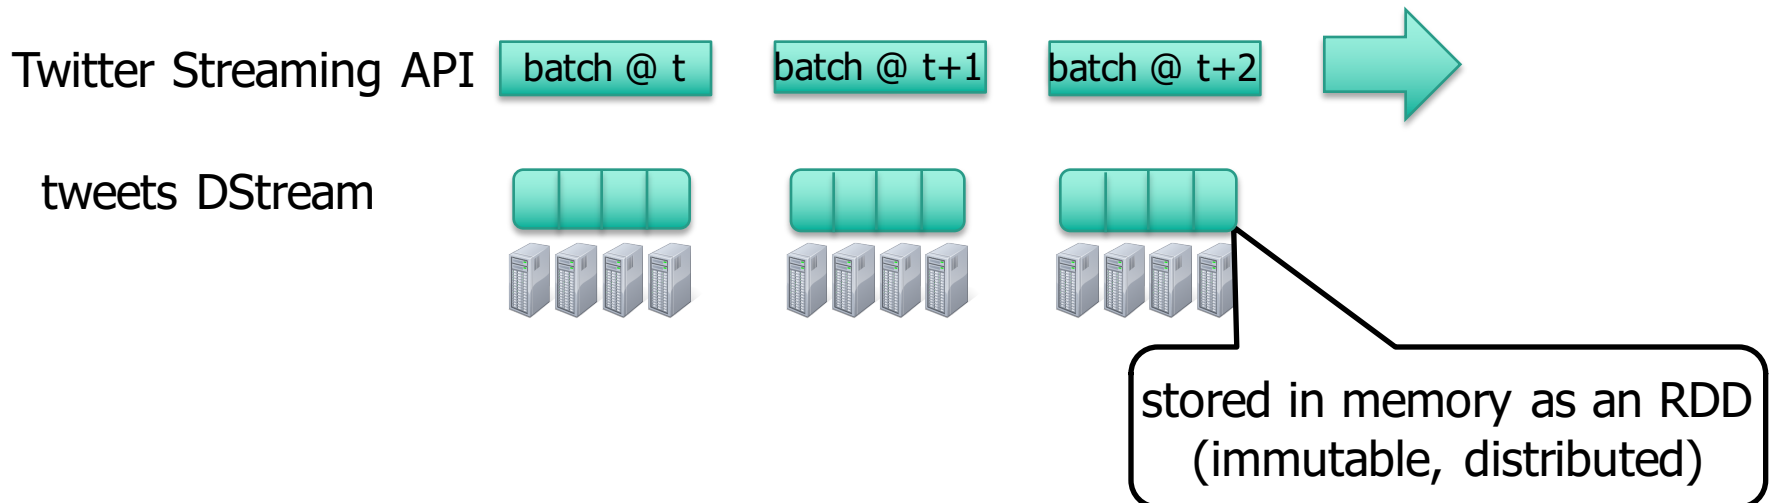
- Batch sizes as low as 1/2 second, live data stream latency of about 1 second
- Potential for combining batch processing and streaming processing in the same system



# Example: Get hashtags from Twitter

```
val tweets = ssc.twitterStream()
```

**DStream**: a sequence of RDDs representing a stream of data

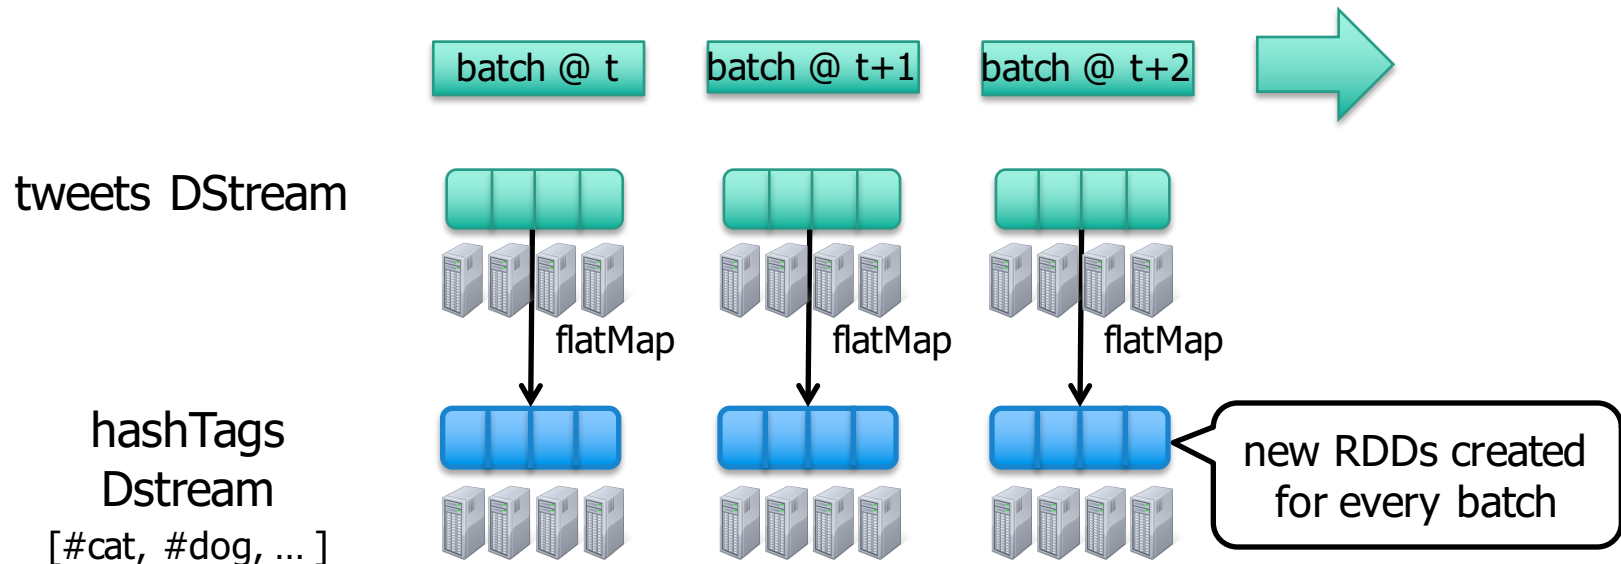


# Example: Get hashtags from Twitter

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))
```

new DStream

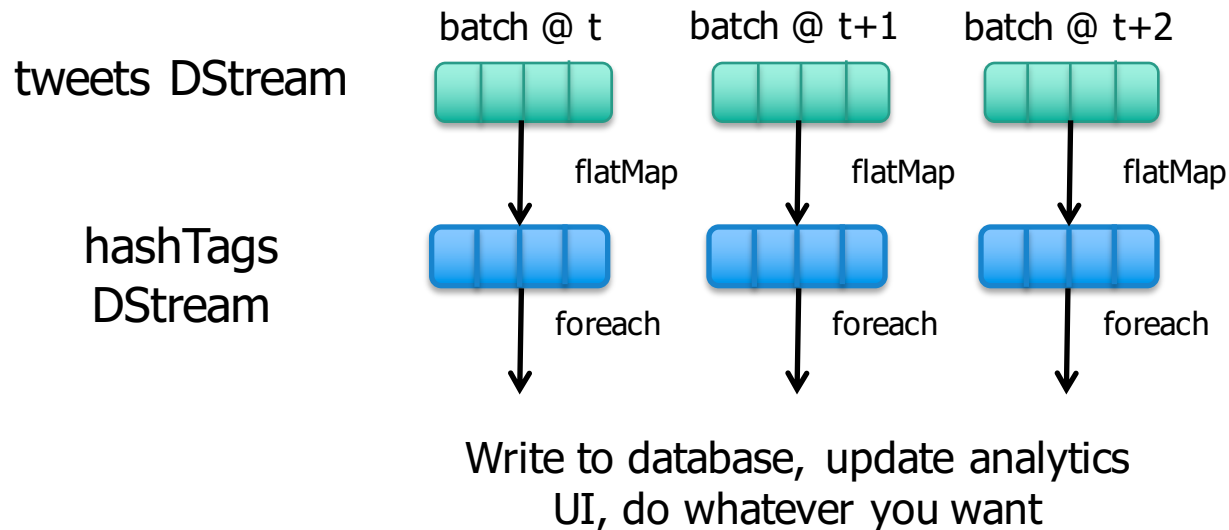
**transformation:** modify data in one DStream to create another DStream



# Example: Get hashtags from Twitter

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))  
hashTags.foreach(hashTagRDD => { ... })
```

**foreach:** do whatever you want with the processed data




# Window-based Transformations

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))  
val tagCounts = hashTags.window(Min(1), Sec(5)).countByValue()
```

sliding window  
operation


window  
length

sliding  
interval



# Fault Tolerance

- RDDs remember the operations that created them
- Batches of input data are replicated in memory for fault-tolerance
- Data lost due to worker failure, can be recomputed from replicated input data





# Streaming Summary

- Stream processing of large amounts of live data is an important requirement
  - Desirable to have both high throughput (100s of MB/s) and low latency ( $\sim$ s)
- Combine the efficiency of in-memory distributed processing of Spark with stream processing model
  - Key is to break down processing in small batches
  - Storm has a second API (Trident) for micro-batch processing
- Also an advantage: use and maintain a single software stack for both processing models

# Stay tuned



<http://www.flickr.com/photos/3dking/2573905313/sizes/l/in/photostream/>

Next time you will learn about:  
**Cloud networks**