

INGI2145 Cloud Computing

Lab 5: Graph processing using Apache Giraph

Muhammad Bilal

Problems with Hadoop MapReduce

- Iterative algorithms are cumbersome and inefficient
- Intermediate results need to persist on HDFS
- Each iteration is separate MR job (scheduled separately)



Apache Giraph




Apache Giraph

- Scalable, fault-tolerant implementation of graph-processing algorithms on Apache Hadoop
- Initially inspired by Google's Pregel framework
- Vertex-centric approach
- Bulk Synchronous Parallelism programming model
- Replaces map reduce with a single mapper only job
- Data partition done once

Apache Giraph

- Giraph algorithms consists of iterative execution of “super-steps”
- Super steps consists of:
 - Message reception
 - Aggregation
 - Vertex/Edge property update
 - Message transmission
- Synchronization phase

Apache Giraph



<http://researcher.ibm.com/researcher/files/us-heq/Large%20Scale%20Graph%20Processing%20with%20Apache%20Giraph.pdf>

Apache Giraph

- Giraph the Vertex Object. $\text{Vertex}\langle I, V, E, M \rangle$
 - I is Vertex ID type
 - V is Vertex Value type
 - E is Edge Value type
 - M is Message data type

Example: N Hop Neighbor count

```
// Initialization in superstep 0
if (getSuperstep() == 0) {
    vertex.setValue(new VertexDataStructure(0l,new HashSet<Long>()));
    for (Edge<LongWritable, FloatWritable> edge : vertex.getEdges()) {
        Message m = new Message(vertex.getId().get(),3,0l);
        sendMessage(edge.getTargetVertexId(), m);
        LOG.debug("Sending message "+ m + " to " + edge.getTargetVertexId());
    }
}
```


Example: N Hop Neighbor count

```
// Process messages and change state
Long nodeCount = vertex.getValue().vertexValue.get();
for (Message message : messages) {
    if (message.srcId.equals(vertex.getId())) {
        LOG.debug("Reply received" + message);
        if (!received.contains(message.senderId.get())) {
            received.add(message.senderId.get());
            LOG.debug(received + " " + vertex.getId());
            nodeCount++;
        }
    } else if (message.hopCount.get() > 0) {
        sendMessage(message.srcId, new Message(message.srcId.get(),
            message.hopCount.get() - 1, vertex.getId().get()));
        for (Edge<LongWritable, FloatWritable> edge : vertex.getEdges()) {
            sendMessage(edge.getTargetVertexId(), new Message(message.srcId.get(),
                message.hopCount.get() - 1, vertex.getId().get()));
        }
    }
}

VertexDataStructure vDS = new VertexDataStructure(nodeCount, received);
vertex.setValue(vDS);
vertex.voteToHalt();
```

Example: N Hop Neighbor count

```
// Process messages and change state
Long nodeCount = vertex.getValue().vertexValue.get();
for (Message message : messages) {
    if(message.srcId.equals(vertex.getId())){
        LOG.debug("Reply received" + message);
        if(!received.contains(message.senderId.get())){
            received.add(message.senderId.get());
            LOG.debug(received + " " +vertex.getId() );
            nodeCount++;
        }
    }else if(message.hopCount.get()>0){
        sendMessage(message.srcId,new Message(message.srcId.get() ,
            message.hopCount.get()-1,vertex.getId().get()));
        for (Edge<LongWritable, FloatWritable> edge : vertex.getEdges()) {
            sendMessage(edge.getTargetVertexId() , new Message(message.srcId.get() ,
                message.hopCount.get()-1,vertex.getId().get()));
        }
    }
}

VertexDataStructure vDS = new VertexDataStructure(nodeCount,received);
vertex.setValue(vDS);
vertex.voteToHalt();
```

Example: N Hop Neighbor count

```
// Process messages and change state
Long nodeCount = vertex.getValue().vertexValue.get();
for (Message message : messages) {
    if(message.srcId.equals(vertex.getId())){
        LOG.debug("Reply received" + message);
        if(!received.contains(message.senderId.get())){
            received.add(message.senderId.get());
            LOG.debug(received + " " +vertex.getId() );
            nodeCount++;
        }
    }else if(message.hopCount.get()>0){
        sendMessage(message.srcId,new Message(message.srcId.get(),
            message.hopCount.get()-1,vertex.getId().get()));
        for (Edge<LongWritable, FloatWritable> edge : vertex.getEdges()) {
            sendMessage(edge.getTargetVertexId(), new Message(message.srcId.get(),
                message.hopCount.get()-1,vertex.getId().get()));
        }
    }
}

VertexDataStructure vDS = new VertexDataStructure(nodeCount,received);
vertex.setValue(vDS);
vertex.voteToHalt();
```



Let's take a look at other parts



Lab Exercise

Shortest Path Computation

Page Rank



Up Next

- Hands-on with Apache Spark

Apache Giraph

- MasterCompute class:
 - Master's compute always runs before the slaves'
 - aggregators are registered here
- Worker Context:
 - Allows execution of user code on a per-worker basis
 - There's one WorkerContext per worker
 - Methods for pre/post superstep operations