

Spark

*Fast, Interactive, Language-Integrated
Cluster Computing*


Muhammad Bilal

*Adapted from ampLabs introductory Spark slides


Why go Beyond MapReduce?

Complex jobs and interactive queries both need one thing that MapReduce lacks:

Efficient primitives for **data sharing**



Iterative algorithm



Interactive data mining

Why go Beyond MapReduce?

Complex jobs and interactive queries both need one thing that MapReduce lacks:

Efficient primitives for **data sharing**



In MapReduce, the only way to share data across jobs is stable storage (e.g. HDFS) -> **slow!**

Iterative algorithm

Interactive data mining

Spark - Features



Extends MapReduce model to better support two common classes of analytics apps:

- » **Iterative** algorithms (machine learning, graphs)
- » **Interactive** data mining

Enhance programmability:



- » Integrate into Scala programming language
- » Allow interactive use from Scala interpreter

Examples



I/O and serialization can take **90%** of the time

Goal: In-Memory Data Sharing



10-100× faster than network and disk

Solution: Resilient Distributed Datasets (RDDs)

Distributed collections of objects that can be stored in memory for fast reuse

Automatically recover lost data on failure

Support a wide range of applications


RDDs maintain **lineage** information that can be used to reconstruct lost partitions

RDD Operations

Transformations (define a new RDD)	map filter sample groupByKey reduceByKey sortByKey	flatMap union join cogroup cross mapValues
Actions (return a result to driver program)	collect reduce count save lookupKey take	

Discretized Streams

- A streaming computation as a series of very small, deterministic batch jobs



Discretized Streams

- **Faults/Stragglers recovery without replication**

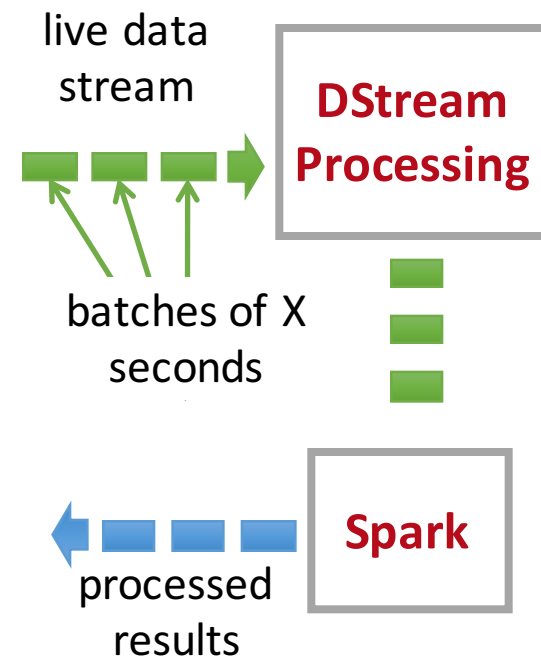
- State between batches kept in memory
- Deterministic operations → fault-tolerant

- **Second-scale performance**

- Try to make batch size as small as possible
- Smaller batch size → lower end-to-end latency

- **A rich set of operations**

- Combined with historical datasets
- Interactive queries



DStream Operations

- Dstream Object: A stream of RDDs
- Transformations
 - Map, filter, groupBy, reduceBy, sort, join...
 - **Windowing**
 - **Incremental aggregation**
- Output operator
 - Save RDDs to outside systems(screen, external storage...
)

Windowing

- Count frequency of words received in last 5 seconds

`words = createNetworkStream("http://...")`


`ones = words.map(w => (w, 1))`

`freqs_5s = ones.reduceByKeyAndWindow(_ + _, Seconds(5), Seconds(1))`


DStream

Transformation


Sliding Window
Ops



Incremental aggregation



Aggregation Function
`freqs = ones.reduceByKeyAndWindow
 (_, +, Seconds(5), Seconds(1))`



Invertible aggregation Function
`freqs =
 ones.reduceByKeyAndWindow
 (_, +, -, Seconds(5), Seconds(1))`

Tutorial Exercises

- Refer to the Lab 6 handout
 - Machine learning and Streaming exercises
- You can reference Spark's programming guides for extra help:
 - <https://spark.apache.org/docs/1.4.1/programming-guide.html>